

RTAI: Embedded Linux VS Legacy RTOS

Gabriele Manduchi
Antonio Barbalace

Introduction

Linux VS RTOS

Linux VS RTOS


Linux (or Linux Embedded)

- General purpose OS;
- MMU;
- Kernel call through syscall (software interrupt).

RTOS like VxWorks, pSOS, Nucleus

- Special Purpose OS;
- Flat Memory model;
- Kernel call through well defined function calls.

Linux VS RTOS




RTOS flat memory model

MAJOR DRAWBACK:
No memory protection

Linux VS RTOS

Linux memory model



QUESTIONS:

RT in kernel space?

...maintain legacy approach

RT in user space?

...exploit memory protection

Linux

From the User Space
As a soft RTOS

Linux

Real-Time constraints in:

- Multimedia application;
- VoIP;
- Audio and videos streaming;
- ...

Linux is a soft real-time system.

Linux

Features of an RTOS:

- Multitasking/multithreading
- Priorities
- Priority inheritance
- Preemption
- Interprocess communication and synchronization
- Dynamic memory allocation

Linux

Why Linux isn't (or wasn't soft) Real Time

- Coarse Grained Synchronization
- Paging
- “Fairness” in Scheduling
- Request Reordering
- Batching

Linux User Space Example

Interrupt flow in Linux for a task waiting for an I/O from disk:

- The I/O is complete. The device raises an interrupt. This cause the block device driver's ISR to run.
- The ISR checks the driver wait queue and finds a task waiting for I/O. It then calls one of the wake-up family functions. The function removes the task from the wait queue and adds it to the scheduler run queue.
- The kernel then calls the function `schedule()` when it gets to a point where scheduling is allowed.
- Finally `schedule()` finds the next suitable candidate for running. The kernel context switches to our task if it has sufficient high priority to get scheduled.

Linux User Space Example

Kernel response time:

- Interrupt latency
- ISR duration
- Scheduler latency
- Scheduler duration

...how was (or will be) improved to have a deterministic and possibly low latency response

Linux User Space Example

Interrupt Latency

PROBLEMS:

- Disabling all interrupts for a long time (removal of all global cli in kernel 2.5);
- Registering a fast interrupt handler by improperly written device driver (mostly a programmer error).

Linux User Space Example

ISR duration

SOLUTION:

Interrupt handler in two stages

BUT:

Soft irq are not deterministic

Linux User Space Example

Scheduler latency

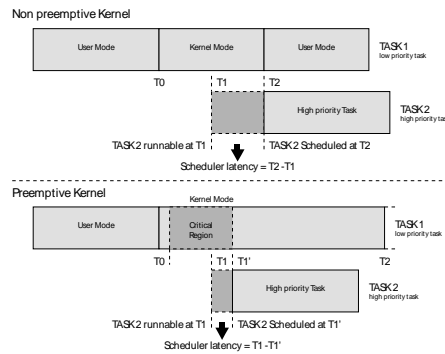
PROBLEMS:

- Nonpreemptive nature of the kernel
- Interrupt disable times

SOLUTIONS:

- Kernel Preemption – by Robert Love
- Low-Latency Patches – by Ingo Molnar and Andrew Morton

Linux User Space Example



Scheduler latency in preemptable and nonpreemptable kernels

Linux User Space Example

Scheduler duration

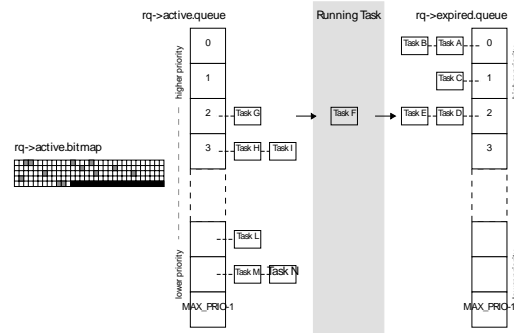
PROBLEMS:

- Non deterministic and slowly goodness and recalculation loop
- Different context switch time between processes and threads (no solution)

SOLUTION:

O(1) Scheduler – Ingo Molnar

Linux User Space Example



Scheduler latency in preemptable and nonpreemptable kernels

Linux

Not only User Space
Approaches for an hard RTOS

Linux hard RTOS

Approaches:

- Mono kernel (changing the source code along different Vanilla kernel)
- Dual kernel (only small changes to the Linux code)

Linux hard RTOS

Dual kernel

(most free)


- RTAI/Xenomai
- RTLinuxFree
- ...

Mono kernel

(most commercials)

- BlueCat Linux
- Cadenux
- Metrowerks
- MontaVista Linux
- RTLinuxPro
- TimeSys Linux
- ...

Linux hard RTOS



Dual kernel approach: two OSs sharing the same hardware, how?

ADEOS

Adaptive Domain Environment for
Operating Systems

ADEOS

Nano kernel

Open source


NEW NAME:

ADEOS/pipe

CHANGES:

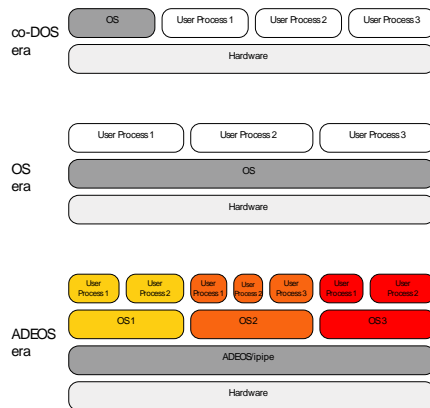
- Functions name
- Various improvements

ADEOS




Dual Kernel with ADEOS/pipe nano kernel

ADEOS



Many different OS (or multiple instance of the same OS) running on the same machine sharing the same hardware.

ADEOS



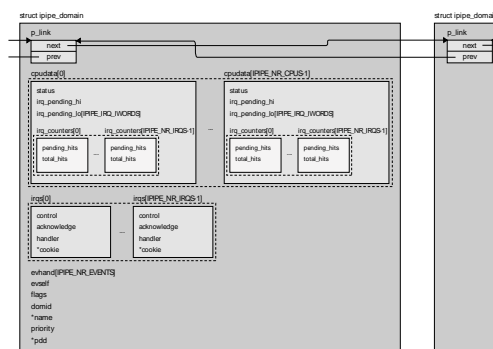
ADEOS/ipipe paths

ADEOS

Basics:


- Domains
- Pipeline
- Interrupts
- Events

ADEOS



ADEOS/ipipe Domain descriptor

ADEOS



ADEOS/pipe interrupts/events pipeline
(on interrupt)

ADEOS

Interrupts

PROPERTIES:

- Asynchronous
- Raised in Hardware
- Maskable (a domain could be stalled)
- Propagated down the pipeline
- Logged (Optimistic interrupt protection scheme)

ADEOS

Events

PROPERTIES:

- Synchronous
- Raised mostly for software cause
- Non maskable
- Propagated from the source domain up the pipeline
- Immediate dispatching

ADEOS

```
struct ipipe_domain domain_desc;
void domain_entry (int iflag) {
    if (iflag) {
        ipipe_virtualize_irq(...);
        ipipe_virtualize_irq(...);
        ipipe_virtualize_irq(...);
        ipipe_catch_event(...);
    }
    for (;;)
        ipipe_suspend_domain();
}

int init_module (void) {
    struct ipipe_domain_attr attr;
    ipipe_init_attr (&attr);
    attr.name = "MyDomain";
    attr.priority = IPIPE_ROOT_PRIO + 100; /* Precede Linux in the pipeline */
    attr.entry = &domain_entry;
    return ipipe_register_domain(&domain_desc,&attr);
}

void cleanup_module (void) {
    return ipipe_unregister_domain(&domain_desc);
}
```

ADEOS/ipipe domain creation example

ADEOS

References:

- Karim Yagmour et others (OperSys Inc)
- www.adeos.org

RTAI

Real Time Application Interface
by Dipartimento di Ingegneria
Aerospaziale dell'Università di Milano
(DIAPM)

RTAI

History (1/2):


- In the first it was PCDOS-DIAPM-RTOS (16bits real mode);
- Than it was implemented in 32 bits real/protected mode;
- Approaching Linux (2.0.25);
- NMT-RTL appear and goes on DIAPM-RTL variant was born
- RTHAL-RTAI patch for 2.2.xx kernel ready (beginning 1999)

RTAI


History (2/2):

- April 1999 first version of RTAI for UP and SMP;
- Jan 2001 NMT-RTL now RTLinux was patented;
- Jun 2002 ADEOS was firstly released;
- ADEOS substitutes RTHAL and from now no patent problems;
- LXRT scheduler was introduced;
- More than one branch of RTAI;
- October 2005 Fusion branch become a stand-alone distribution Xenomai;

RTAI VS Xenomai



RTAI immediate dispatching



Xenomai

RTAI pipeline (no immediate dispatching)

RTAI VS Xenomai

RTAI

Currently ported on:

- ARM
- I386, ia64 (SMP, MUP)
- PowerPC

Run:

- Kernel
- User

Main goal :

- Fast dispatching

API interface:

- RTAI
- Posix 1003.1b
- RTDM

Xenomai

Currently ported on:

- ARM
- i386, ia64 (SMP)
- PowerPC, PowerPC64 (SMP)
- Blackfin

Run:

- Kernel
- User

Main goal :

- Portability

API interface:

- Xenomai native
- Posix 1003.1b
- RTAI
- VxWorks
- pSos+
- VRTX
- ulTRON
- RTDM

RTAI VS Xenomai

From our tests:

LESS INTEGRATION PROBLEMS:

Xenomai

BETTER PERFORMANCE:

RTAI (immediate dispatching)

EASY APPLICATION PORTING:

Xenomai

DEBUGGING SUPPORT:

Both

CLEAR, AND WELL ORGANIZED SOURCE CODE:

Xenomai


RTAI

Immediate dispatching

RTAI

Basics:

- ADEOS/RTHAL
- RTAI Scheduler
- RTAI IPC



RTAI

ADEOS/RTHAL

PROVIDES:

- Immediate Interrupts dispatching
- Interrupt registering/deregistering functions
- PIC management routines
- Hardware Timers (8254/APIC) management routines
- registering/deregistering timers handler functions
- Trap handlers registering/deregistering functions
- Provides `rt_printk/rt_synch_printk`
- Syscall dispatching

RTAI

ADEOS/RTHAL interrupts

Pure ADEOS Xenomai

Raised an External interrupt n

CPU decode phase

CPU call IDT[n] entry

```
linux/arch/i386/kernel/entry.S:404 interrupt[n]()
linux/arch/i386/kernel/entry.S:475 common_interrupt()
linux/arch/i386/kernel/ipipe-root.c: __ipipe_handle_irq(struct pt_regs
regs)
detect head Domain (Xenomai) is wired
linux/kernel/ipipe/core.c: __ipipe_dispatch_wired(struct ipipe_domain
*head, unsigned irq)
call Xenomai interrupt handler (registered with rthal_irq_request)
```

RTHAL ADEOS modified patch

Raised an External interrupt n

CPU decode phase

CPU call IDT[n] entry

```
linux/arch/i386/kernel/entry.S:404 interrupt[n]()
linux/arch/i386/kernel/entry.S:475 common_interrupt()
rtai/base/arch/i386/hal/hal.immed:1474 rtai_hirq_dispatcher(struct
pt_regs regs)
call RTAI interrupt handler (register with rt_request_irq)
```

RTAI

ADEOS/RTHAL interrupts And Linux interrupts?

Before RTAI

Raised an External interrupt n

CPU decode phase

CPU call IDT[n] entry

```
linux/arch/i386/kernel/entry.S:404 interrupt[n]()
linux/arch/i386/kernel/entry.S:416 common_interrupt()
linux/arch/i386/kernel/irq.c:54 do_IRQ(struct pt_regs *regs)
linux/kernel/irq/handle.c:118 __do_IRQ(unsigned int irq, struct pt_regs *regs)
linux/kernel/irq/handle.c:79 handle_IRQ_event(unsigned int irq, struct pt_regs *regs, struct irqaction
*action)
loop through the actions and call the interrupt handlers
```

With RTAI

Raised an External interrupt n

CPU decode phase

CPU call IDT[n] entry

```
linux/arch/i386/kernel/entry.S:404 interrupt[n]()
linux/arch/i386/kernel/entry.S:475 common_interrupt()
rtai/base/arch/i386/hal/hal.c:1474 rtai_hirq_dispatcher(struct pt_regs regs)
call RTAI interrupt handler
rtai/base/include/asm-i386/rtai_hal.h:275 hal_pend_domain_uncond(irq, domain, cpuid)
rtai/base/include/asm-i386/rtai_hal.h:285 hal_fast_flush_pipeline(cpuid)
linux/arch/i386/kernel/ipipe-core.c: __ipipe_sync_stage(unsigned long syncmask)
for every interrupt pended in the log do
call handler
linux/arch/i386/kernel/irq.c:54 do_IRQ(struct pt_regs *regs)
linux/kernel/irq/handle.c:118 __do_IRQ(unsigned int irq, struct pt_regs *regs)
linux/kernel/irq/handle.c:79 handle_IRQ_event(unsigned int irq, struct pt_regs *regs, struct irqaction
*action)
loop through the actions and call the interrupt handlers
```

RTAI

8254

Native RTAI timer, it supports at least two modes that are used:

- Periodic Mode
- Oneshot Mode

This timer is used in:

- UP
- SMP

Counter 0 shared with Linux.

ADEOS/RTHAL timers

APIC

In the beginning was used only with multiprocessor hardware.


Same modes:

- Periodic Mode
- Oneshot Mode

Used in:

- SMP
- MUP

RTAI



Periodic Mode

Automatic timer register reload

Oneshot Mode

Compulsory timer register update

RTAI

```
bstatic int irq_handler(int irq, void *cookie) {
    // do something in IRQ handler

    rt_enable_irq(vmechip_irq);
    return IRQ_HANDLED;
}

static int VGD4_init(void) {
    result = rt_request_irq(vmechip_irq, VGD4_irqhandler, 0, 1);
    if (result != 0) {
        printk(KERN_ERR"VG4 driver: can't get assigned pci irq vector %02X\n", vmechip_irq);
        return -1;
    } else {
        rt_enable_irq(vmechip_irq);
    }

    return 0;
}

static void VGD4_exit(void) {
    rt_disable_irq(vmechip_irq);
    rt_release_irq(vmechip_irq);

    return;
}
```

ADEOS/RTHAL interrupt handler registration
example

RTAI

Scheduler

PROVIDES:

- Task creation/destruction routine
- Task suspend/yield/sleep routine
- Task Timer related function
- Scheduling disciplines/policy
- LXRT Linux (user space) Real-Time capability
- Task interface with IPC mechanisms

RTAI

Scheduler

MAIN ENTITY: task

```
typedef struct rt_task_struct {
    long *stack __attribute__((__aligned__(L1_CACHE_BYTES)));
    int uses_fpu;
    int magic;
    volatile int state, running;
    unsigned long runnable_on_cpu;
    long *stack_bottom;
    volatile int priority;
    int base_priority;
    int policy;
    int sched_lock_priority;
    struct rt_task_struct *prio_passed_to;
    RTIME period;
    RTIME resume_time;
    RTIME periodic_resume_time;
    RTIME yield_time;
    int rr_quantum;
    int rr_remaining;
    int suspend;
    struct rt_queue queue;
    int owdress;
    struct rt_queue *blocked_on;
    struct rt_queue msg_queue;
    int tid; /* trace ID */
    unsigned long msg;
    struct rt_queue ret_queue;
    void (*signal)(void);
    FPU_ENV fpu_reg __attribute__((__aligned__(L1_CACHE_BYTES)));
    struct rt_task_struct *prev;
    struct rt_task_struct *next;
    struct rt_task_struct *tprev;
    struct rt_task_struct *tnext;
    struct rt_task_struct *rprev;
    struct rt_task_struct *rnext;

    /* Appended for calls from LINUX. */
    long *fun_args;
    long *bstack;
    struct task_struct *lnxtsk;
    long long retval;
    char *msg_buf[2];
    long max_msg_size[2];
    char task_name[16];
    void *system_data_ptr;
    struct rt_task_struct *nextp;
    struct rt_task_struct *prevp;

    /* Added to support user specific trap handlers. */
    RT_TRAP_HANDLER task_trap_handler[HAL_NR_FAULTS];

    /* Added from rtai-22. */
    long unblocked;
    void *rt_signals;
    volatile unsigned long pstate;
    unsigned long usp_flags;
    unsigned long usp_flags_mask;
    unsigned long force_soft;
    volatile int is_hard;

    void *trap_handler_data;
    struct rt_task_struct *linux_syscall_server;

    /* For use by watchdog. */
    int resync_frame;

    /* For use by exit handler functions. */
    XHDL *ExitHook;

    RTIME exectime[2];
    struct tcb_t tcb;

    /* Real time heaps. */
    struct rt_heap_t heap[2];

    volatile int scheduler;

#ifdef CONFIG_RTAI_LONG_TIMED_LIST
    rb_root_t rbr;
    rb_node_t rbn;
#endif
    struct rt_queue resq;
} RT_TASK __attribute__((__aligned__(L1_CACHE_BYTES)));
```

RTAI

task **RT_TASK**

3 double linked lists:

- fields prev/next **task list**
- fields rprev/rnext **ready list**
- fields fprev/fnext **timed list**

A list head for every CPU

rt_smp_linux_task[cupid]

RTAI

task **RT_TASK**
task list

INSERTION POLICY:

Add every new created task

LIST ORDER:

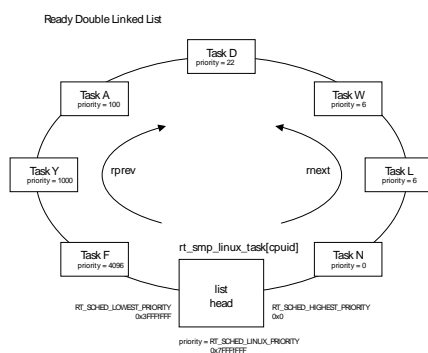
FIFO

NOTE:

Tail task next = 0

RTAI

task **RT_TASK**
ready list



INSERTION POLICY:

Add ready task

REMOVE POLICY:

First task after the head

LIST ORDER:

Sorted on `rt_task.priority`

MOST PRIORITARY:

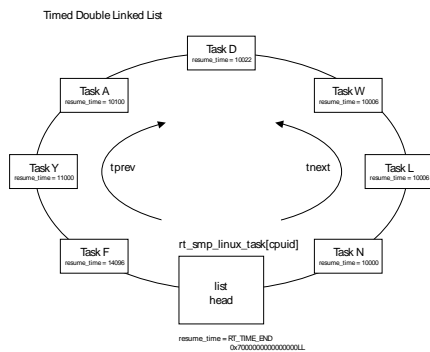
0

LESS PRIORITARY:

`RT_SCHED_LINUXPRIO`

RTAI

task **RT_TASK**
timed list



INSERTION POLICY:

Add timed task

REMOVE POLICY:

First n deadlined tasks after the head

LIST ORDER:

Sorted on
rt_task.resume_time

MOST PRIORITY:

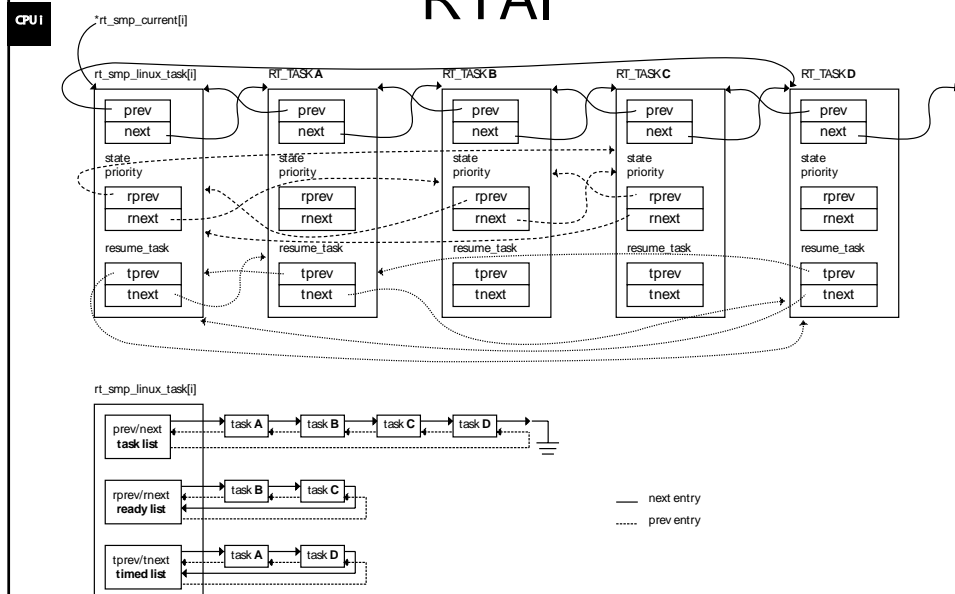
0

LESS PRIORITY:

MAX_TIME

RTAI

for every CPU i



RTAI

Scheduler

Scheduling policy:

- RMS
- EDF
- FIFO, **default scheduler**
- RR

RTAI

Scheduler

Scheduling point:

- on timer interrupt
- on ISR exit
- on task self-suspending

Self-suspending examples:

- Yield
- Suspend
- Sleep

RTAI

Scheduler

LXRT

Scheduler kernel

Scheduler user

RTAI

```
void task_body (long cookie) {
    int i, l;
    while(1){
        for(i=0; i<100; i++) {
            l += i*(1+i);
        }
        printk(KERN_ERR"kernel_task executed\n");
        rt_task_wait_period();
    }
}

int init_module (void) {
    int err;
    err = rt_task_init(&task_desc, task_body, 0, TASK_STKSZ, TASK_PRIO, 0, 0);
    if (err != 0) {
        printk(KERN_ERR"error module loading\n");
        return -1;
    }
    rt_set_oneshot_mode();
    start_rt_timer(0);
    rt_task_make_periodic_relative_ns(&task_desc, 100000000, 1000000000000);
    return 0;
}

void cleanup_module(void) {
    stop_rt_timer();
    rt_task_delete(&task_desc);
}
```

Scheduler task creation example

RTAI

IPCs

PROVIDES:

- Bits (a sort of signals)
- Fifo
- Message
- Mailbox
- Message queue
- NetRPC (now require RTNet)
- Semaphore
- Shared memory

RTAI

RTAI References:

Paolo Mantegazza, DIAPM

www.rtai.org

Xenomai References:

www.xenomai.org